

# Performance Analysis in C++ with `google/benchmark`

As you optimize code, you will begin to develop intuitions about what C++ code is fast and what is slow

☞ `std::array` is faster than `std::vector`

☞ `if` checks break the pipeline

☞ `std::vector` is faster than `std::list`

☞ `std::vector::reserve` is faster than  
`std::vector::push_back`

But tomorrow a smart compiler writer could make all this stuff false.

# What does a performance benchmark need to do?

- ☞ Call the function repeatedly until statistical confidence is gained about its runtime, and no longer
- ☞ Not get optimized out by the compiler
- ☞ Test multiple inputs to determine asymptotic scaling
- ☞ Be usable

google/benchmark fits the bill

# Installation

```
$ git clone https://github.com/google/benchmark.git
$ mkdir build_bm; cd build_bm
build_bm$ cmake -DCMAKE_BUILD_TYPE=Release ../benchmark
build_bm$ make -j`nproc`
build_bm$ make test
build_bm$ sudo make install
```

# google/benchmark mwe

```
#include <cmath>
#include <benchmark/benchmark.h>

static void BM_Pow(benchmark::State& state)
{
    while (state.KeepRunning())
    {
        auto y = std::pow(1.2, 1.2);
    }
}

BENCHMARK(BM_Pow);

BENCHMARK_MAIN();
```

# Build sequence:

```
all: run_benchmarks.x run_benchmarks.s
```

```
run_benchmarks.x: run_benchmarks.o  
    $(CXX) -o $@ $< -lbenchmark -pthread
```

```
run_benchmarks.o: run_benchmarks.cpp  
    $(CXX) -std=c++14 -O3 -c $< -o $@
```

```
run_benchmarks.s: run_benchmarks.cpp  
    $(CXX) $(CXXFLAGS) -S -masm=intel $<
```

```
clean:  
    rm -f *.x *.s *.o
```

# Run a google/benchmark

```
$ ./run_benchmarks.x
```

```
Run on (4 X 1000 MHz CPU s)
```

```
2016-06-23 17:58:41
```

Benchmark	Time	CPU	Iterations
BM_Pow	3 ns	3 ns	264837522



# Stop compiler optimizations

- ☞ 3 ns seems a bit fast for this operation.
- ☞ The compiler might have (correctly) reasoned that the repeated call to `std::pow` is useless, and optimized it out.
- ☞ We can generate the assembly of this function via  
`clang++ -std=c++14 -O3 -S -masm=intel run_benchmarks.cpp`

# Stop compiler optimizations

We can see all function calls in the assembly via

```
$ cat run_benchmarks.s | grep 'call' | awk '{print $2}' | xargs c++filt
benchmark::State::KeepRunning()
benchmark::Initialize(int*, char**)
benchmark::RunSpecifiedBenchmarks()
benchmark::State::ResumeTiming()
benchmark::State::PauseTiming()
__assert_fail
operator new(unsigned long)
benchmark::internal::Benchmark::Benchmark(char const*)
benchmark::internal::RegisterBenchmarkInternal(benchmark::internal::Benchmark*)
operator delete(void*)
_Unwind_Resume
```

`std::pow` isn't one of them!

# Stop compiler optimizations

- ☞ The compiler's goal is to remove all unnecessary operations from your code
- ☞ Your goal is to do unnecessary operations to see how long a function call takes

# Stop compiler optimizations

☞ Benchmarked writes being optimized out is a huge problem.

☞ google/benchmark has created a function to deal with it:

`benchmark::DoNotOptimize`

```
double y;  
while (state.KeepRunning()) {  
    benchmark::DoNotOptimize(y = std::pow(1.2, 1.2));  
}
```

# Stop compiler optimizations:

`benchmark::DoNotOptimize` forces the result to be stored into RAM

```
template <class Tp>
inline BENCHMARK_ALWAYS_INLINE void DoNotOptimize(Tp const& value) {
    asm volatile("" : "+m" (const_cast<Tp&>(value)));
}
```

# Stop compiler optimizations

The purpose of this is to tell the compiler to not optimize out the assignment of `y`.

But `benchmark::DoNotOptimize` can't keep the compiler from evaluating `std::pow(1.2, 1.2)` at compile time.

# Stop compiler optimizations

To keep the compiler from evaluating `std::pow(1.2, 1.2)` at compile time, we simply need to ensure that it doesn't know what values it needs to evaluate.

```
std::random_device rd;
std::mt19937 gen(rd());
std::uniform_real_distribution<double> dis(1, 10);
auto s = dis(gen);
auto t = dis(gen);
double y;
while (state.KeepRunning())
{
    benchmark::DoNotOptimize(y = std::pow(s, t));
}
```

# Stop compiler optimizations

Even then we might still have to play tricks on the compiler. One of my favorites: Write the result to `/dev/null` outside the loop:

```
double y;  
while (state.KeepRunning()) {  
    benchmark::DoNotOptimize(y = std::pow(s, t));  
}  
std::ostream cnull(0);  
cnull << y;
```



# Full boilerplate

```
#include <cmath>
#include <ostream>
#include <random>
#include <benchmark/benchmark.h>

static void BM_Pow(benchmark::State& state) {
    std::random_device rd;
    std::mt19937 gen(rd());
    std::uniform_real_distribution<double> dis(1, 10);
    auto s = dis(gen);
    auto t = dis(gen);
    double y;
    while (state.KeepRunning()) {
        benchmark::DoNotOptimize(y = std::pow(s, t));
    }
    std::ostream cnull(0);
    cnull << y;
}

BENCHMARK(BM_Pow);
BENCHMARK_MAIN();
```

# Stop compiler optimizations

Now our timings are more in line with our expectations:

```
$ ./run_benchmarks.x
```

```
Run on (1 X 2300 MHz CPU )
```

```
2016-06-24 20:11:40
```

Benchmark	Time	CPU	Iterations
BM_Pow	80 ns	80 ns	9210526

# Templated Benchmarks

It's often useful to find out how fast your algorithm is in float, double, and long double precision. Google benchmark supports templates without too much code duplication

# Templated Benchmarks

```
template<typename Real>
static void BM_PowTemplate(benchmark::State& state) {
    std::random_device rd;
    std::mt19937 gen(rd());
    std::uniform_real_distribution<Real> dis(1, 10);
    auto s = dis(gen);
    auto t = dis(gen);
    Real y;
    while (state.KeepRunning()) {
        benchmark::DoNotOptimize(y = std::pow(s, t));
    }
    std::ostream cnul(nullptr);
    cnul << y;
}
BENCHMARK_TEMPLATE(BM_PowTemplate, float);
BENCHMARK_TEMPLATE(BM_PowTemplate, double);
BENCHMARK_TEMPLATE(BM_PowTemplate, long double);
```

# Templated Benchmarks

The results are sometimes surprising; for instance double is found to be faster than float:

```
$ ./run_benchmarks.x
```

```
Run on (1 X 2300 MHz CPU )
```

```
2016-06-25 00:07:26
```

Benchmark	Time		CPU Iterations	
-----				
BM_PowTemplate<float>	136	ns	127	ns 5468750
BM_PowTemplate<double>	95	ns	94	ns 7000000
BM_PowTemplate<long double>	404	ns	403	ns 1699029

# Determine asymptotic scaling

Recursive Fibonacci numbers:

```
uint64_t fibr(uint64_t n)
{
    if (n == 0)
        return 0;

    if (n == 1)
        return 1;

    return fibr(n-1)+fibr(n-2);
}
```

# Determine asymptotic scaling

```
static void BM_FibRecursive(benchmark::State& state)
{
    uint64_t y;
    while (state.KeepRunning())
    {
        benchmark::DoNotOptimize(y = fibr(state.range_x()));
    }
    std::ostream cnull(nullptr);
    cnull << y;
}
BENCHMARK(BM_FibRecursive)->RangeMultiplier(2)->Range(1, 1<<5);
```

# Determine asymptotic scaling

```
BENCHMARK(BM_FibRecursive)->RangeMultiplier(2)->Range(1, 1<<5)
```

will request a benchmark with `state.range_x()` taking values of `[1, 2, 4, 8, 16, 32]`.



# Determine asymptotic scaling

Run on (1 X 2300 MHz CPU )

2016-06-25 00:38:14

Benchmark	Time		CPU Iterations	
-----				
BM_FibRecursive/1	7	ns	7 ns	83333333
BM_FibRecursive/2	15	ns	15 ns	72916667
BM_FibRecursive/4	37	ns	37 ns	17156863
BM_FibRecursive/8	268	ns	268 ns	2868852
BM_FibRecursive/16	13420	ns	13392 ns	64815
BM_FibRecursive/32	24372253	ns	24320000 ns	25

# Determine asymptotic scaling

Can we empirically determine the asymptotic complexity of the recursive Fibonacci number calculation?

Pretty much . . . !

# Determine asymptotic scaling

google/benchmark will try to figure out the asymptotic scaling, if add a `Complexity()` call:

```
BENCHMARK(BM_FibRecursive)
  ->RangeMultiplier(2)
    ->Range(1, 1<<5)
      ->Complexity();
```

# Result

```
$ ./run_benchmarks.x
```

BM_FibRecursive/1	9 ns	9 ns	72916667
BM_FibRecursive/2	19 ns	19 ns	44871795
BM_FibRecursive/4	42 ns	43 ns	14112903
BM_FibRecursive/8	270 ns	268 ns	2611940
BM_FibRecursive/16	11305 ns	11264 ns	54687
BM_FibRecursive/32	27569038 ns	27555556 ns	27
BM_FibRecursive_Big0	828.24 N^3	827.83 N^3	
BM_FibRecursive_RMS	31 %	31 %	

It erroneously labels the algorithm as cubic; though we can see the fit is not tight.

# Pass a lambda to Complexity()

If google/benchmark only tries to fit to the most common complexity classes. You are free to specify the asymptotic complexity yourself, and have google/benchmark determine goodness of fit.

The complexity of the recursive Fibonacci algorithm is  $\phi^n$ , where  $\phi := (1 + \sqrt{5})/2$  is the golden ratio, so let's use  $\phi^n$  as a lambda . . .

# Pass a lambda to Complexity()

Syntax:

```
BENCHMARK(BM_FibRecursive)
  ->RangeMultiplier(2)
  ->Range(1, 1<<5)
  ->Complexity([](int n) {return std::pow((1+std::sqrt(5))/2, n);});
```

google/benchmark nails it:

BM_FibRecursive/1	9 ns	9 ns	79545455
BM_FibRecursive/2	19 ns	19 ns	44871795
BM_FibRecursive/4	37 ns	37 ns	20588235
BM_FibRecursive/8	228 ns	228 ns	3125000
BM_FibRecursive/16	9944 ns	9943 ns	60345
BM_FibRecursive/32	23910141 ns	23866667 ns	30
BM_FibRecursive_Big0	4.91 f(N)	4.90 f(N)	
BM_FibRecursive_RMS	0 %	0 %	

(Note: All lambdas are denoted as  $f(N)$ , so you have to know what you wrote in your source code to understand the scaling.)

# Other tricks:

Standard complexity classes don't need to be passed as lambdas:

```
Complexity(benchmark::o1);
```

```
Complexity(benchmark::oLogN);
```

```
Complexity(benchmark::oN);
```

```
Complexity(benchmark::oNLogN);
```

```
Complexity(benchmark::oNSquared);
```

```
Complexity(benchmark::oNCubed);
```



# Benchmarking data transfer

Use `state.SetBytesProcessed`<sup>1</sup>:

```
static void BM_memcpy(benchmark::State& state) {  
    char* src = new char[state.range_x()]; char* dst = new char[state.range_x()];  
    memset(src, 'x', state.range_x());  
    while (state.KeepRunning()) {  
        memcpy(dst, src, state.range_x());  
    }  
    state.SetBytesProcessed(int64_t(state.iterations()) *  
                           int64_t(state.range_x()));  
    delete[] src; delete[] dst;  
}  
BENCHMARK(BM_memcpy) -> Range(8, 8<<10);
```

---

<sup>1</sup> Taken directly from the [google/benchmark](https://github.com/google/benchmark) docs.

# Benchmarking data transfer

`state.SetBytesProcessed` adds another column to the output:

Run on (1 X 2300 MHz CPU )

2016-06-25 19:25:23

Benchmark	Time		CPU		Iterations	
-----						
BM_memcpy/8	9	ns	9	ns	87500000	883.032MB/s
BM_memcpy/16	9	ns	9	ns	79545455	1.70305GB/s
BM_memcpy/32	9	ns	8	ns	79545455	3.50686GB/s
BM_memcpy/64	11	ns	11	ns	62500000	5.35243GB/s
BM_memcpy/128	13	ns	13	ns	53030303	8.97969GB/s
BM_memcpy/256	16	ns	16	ns	44871795	15.1964GB/s
BM_memcpy/512	19	ns	20	ns	36458333	24.4167GB/s
BM_memcpy/1024	25	ns	25	ns	28225806	38.6756GB/s
BM_memcpy/2k	50	ns	50	ns	10000000	38.147GB/s
BM_memcpy/4k	122	ns	122	ns	5833333	31.2534GB/s
BM_memcpy/8k	220	ns	220	ns	3240741	34.726GB/s

# Run a subset of the benchmarks

```
$ ./run_benchmarks.x --benchmark_filter=BM_memcpy/32
```

```
Run on (1 X 2300 MHz CPU )
```

```
2016-06-25 19:34:24
```

Benchmark	Time		CPU		Iterations	
-----						
BM_memcpy/32	11	ns	11	ns	79545455	2.76944GB/s
BM_memcpy/32k	2181	ns	2185	ns	324074	13.9689GB/s
BM_memcpy/32	12	ns	12	ns	54687500	2.46942GB/s
BM_memcpy/32k	1834	ns	1837	ns	357143	16.6145GB/s

# Error bars

- ☞ In general you will get a pretty good idea about the standard deviation of the measurements just by running it a few times.
- ☞ However, if you want error bars, just specify the number of times you want your benchmark repeated:

```
BENCHMARK(BM_Pow)→Repetitions(12);
```

# Error bars

Run on (1 X 2300 MHz CPU )

2016-06-25 19:57:40

Benchmark	Time		CPU Iterations	
BM_Pow/repeats:12	70	ns	70	ns 10294118
BM_Pow/repeats:12	73	ns	73	ns 10294118
BM_Pow/repeats:12	71	ns	71	ns 10294118
BM_Pow/repeats:12	70	ns	70	ns 10294118
BM_Pow/repeats:12	71	ns	71	ns 10294118
BM_Pow/repeats:12	72	ns	72	ns 10294118
BM_Pow/repeats:12	73	ns	73	ns 10294118
BM_Pow/repeats:12	71	ns	71	ns 10294118
BM_Pow/repeats:12	70	ns	70	ns 10294118
BM_Pow/repeats:12	73	ns	73	ns 10294118
BM_Pow/repeats:12	71	ns	71	ns 10294118
BM_Pow/repeats:12	76	ns	75	ns 10294118
BM_Pow/repeats:12_mean	72	ns	72	ns 10294118
BM_Pow/repeats:12_stddev	2	ns	2	ns 0

# google/benchmark gotchas

If you have CPU frequency scaling enabled (think laptops with power saving), then the `Time` column can get inaccurate:

Run on (1 X 2300 MHz CPU )

2016-06-25 19:57:40

Benchmark	Time	CPU Iterations
-----		

Next run:

Run on (1 X 2750 MHz CPU )

2016-06-25 19:57:40

Benchmark	Time	CPU Iterations
-----		

# google/benchmark gotchas

I suspect negative wall times are due to the CPU frequency changing during the course of computation.

Unfortunately, at the time of this writing, this issue is unresolved.