

Performance Analysis in C++

The principle challenge of HPC is *understanding your stack*.

The *first* challenge is *defining correctness*, and asserting on invariants.

The *second* challenge is market competitiveness, which is based principally on usability.

- >> Adaptive numerical algorithms instead of user parameters
- >> Progress reporting
- >> Widely compatible data formats
- >> Curated algorithms, rather than *choice* of algorithms

The *last* challenge is performance analysis, which is today's topic.

As you optimize code, you will begin to develop intuitions about what C++ code is fast and what is slow

>> `std::array` is faster than `std::vector`

>> `float` is faster than `double`

>> `if` checks break the pipeline

>> `std::vector` is faster than `std::list`

>> `std::vector::reserve` is faster than `std::vector::push_back`

But tomorrow a smart compiler writer could make all this stuff false.

What does a performance benchmark need to do?

- >> Call the function repeatedly until statistical confidence is gained about its runtime, and no longer
- >> Not get optimized out by the compiler
- >> Test multiple inputs for signs of scaling problems
- >> Be usable

Use google/benchmark!

The installation is simple:

```
$ git clone https://github.com/google/benchmark.git
$ mkdir build_benchmark; cd build_benchmark
build_benchmark$ cmake -DCMAKE_BUILD_TYPE=Release ../benchmark
build_benchmark$ make -j`nproc`
build_benchmark$ make test
build_benchmark$ sudo make install
```


google/benchmark minimal working example

```
#include <cmath>
#include <benchmark/benchmark.h>

static void BM_Pow(benchmark::State& state)
{
    while (state.KeepRunning())
    {
        auto y = std::pow(1.2, 1.2);
    }
}

BENCHMARK(BM_Pow);

BENCHMARK_MAIN();
```

google/benchmark minimal working example

Build sequence:

```
CXX=clang++
```

```
all: run_benchmarks.x run_benchmarks.s
```

```
run_benchmarks.x: run_benchmarks.o  
    $(CXX) -o $@ $< -lbenchmark -pthread
```

```
run_benchmarks.o: run_benchmarks.cpp  
    $(CXX) -std=c++14 -O3 -c $< -o $@
```

```
run_benchmarks.s: run_benchmarks.cpp  
    $(CXX) $(CPPFLAGS) -S -masm=intel $<
```

```
clean:  
    rm -f *.x *.s *.o
```

google/benchmark minimal working example

Run:

`./run_benchmarks.x`

Run on (4 X 1000 MHz CPU s)

2016-06-23 17:58:41

| Benchmark | Time | CPU | Iterations |
|-----------|------|------|------------|
| ----- | | | |
| BM_Pow | 3 ns | 3 ns | 264837522 |

Stop compiler optimizations

In fact 3 ns seems a bit fast for this operation. The compiler might have (correctly) reasoned that the repeated call to `std::pow` is useless, and optimized it out.

We can generate the assembly of this function via

```
clang++ -std=c++14 -O3 -S -masm=intel run_benchmarks.cpp
```

Stop compiler optimizations

We can see all function calls in the assembly via

```
$ cat run_benchmarks.s | grep 'call' | awk '{print $2}' | xargs c++filt
benchmark::State::KeepRunning()
benchmark::Initialize(int*, char**)
benchmark::RunSpecifiedBenchmarks()
benchmark::State::ResumeTiming()
benchmark::State::PauseTiming()
__assert_fail
operator new(unsigned long)
benchmark::internal::Benchmark::Benchmark(char const*)
benchmark::internal::RegisterBenchmarkInternal(benchmark::internal::Benchmark*)
operator delete(void*)
_Unwind_Resume
```

std::pow isn't one of them!

Stop compiler optimizations

- >> The compiler's goal is to remove all unnecessary operations from your code
- >> Your goal is to do unnecessary operations to see how long a function call takes

Stop compiler optimizations

This problem is so pervasive that `google/benchmark` has created a function to deal with it: `benchmark::DoNoOptimize`:

```
double y
while (state.KeepRunning()) {
    benchmark::DoNotOptimize(y = std::pow(1.2, 1.2));
}
```

Stop compiler optimizations

The purpose of this is to tell the compiler to *not* optimize out the assignment of `y`.

But `benchmark::DoNotOptimize` can't keep the compiler from evaluating `std::pow(1.2, 1.2)` at compile time.

Stop compiler optimizations

To keep the compiler from evaluating `std::pow(1.2, 1.2)` at compile time, we simply need to ensure that it doesn't *know* what values it needs to evaluate. Here's a solution:

```
std::random_device rd;
std::mt19937 gen(rd());
std::uniform_real_distribution<double> dis(1, 10);
auto s = dis(gen);
auto t = dis(gen);
double y;
while (state.KeepRunning())
{
    benchmark::DoNotOptimize(y = std::pow(s, t));
}
```

Stop compiler optimizations:

However, `benchmark::DoNotOptimize` forces the result to be stored into RAM, which takes time over register storage:

```
template <class Tp>
inline BENCHMARK_ALWAYS_INLINE void DoNotOptimize(Tp const& value) {
    asm volatile("" : "+m" (const_cast<Tp&>(value)));
}
```

Stop compiler optimizations

Even then we might still have to play tricks on the compiler. One of my favorites: Write the result to `/dev/null` outside the loop:

```
double y
while (state.KeepRunning()) {
    benchmark::DoNotOptimize(y = std::pow(s, t));
}

std::ostream cnull(0);
cnull << y;
```

Stop compiler optimizations: Full boilerplate

```
#include <cmath>
#include <ostream>
#include <random>
#include <benchmark/benchmark.h>

static void BM_Pow(benchmark::State& state) {
    std::random_device rd;
    std::mt19937 gen(rd());
    std::uniform_real_distribution<double> dis(1, 10);
    auto s = dis(gen);
    auto t = dis(gen);
    double y;
    while (state.KeepRunning()) {
        benchmark::DoNotOptimize(y = std::pow(s, t));
    }
    std::ostream cnull(0);
    cnull << y;
}

BENCHMARK(BM_Pow);
BENCHMARK_MAIN();
```

Stop compiler optimizations

Now our timings are more in line with our expectations:

Run on (1 X 2300 MHz CPU)

2016-06-24 20:11:40

| Benchmark | Time | CPU | Iterations |
|-----------|-------|-------|------------|
| ----- | | | |
| BM_Pow | 80 ns | 80 ns | 9210526 |

Templated Benchmarks

It's often useful to find out how fast your algorithm is in float, double, and long double precision. Google benchmark supports templates without too much code duplication

Templated Benchmarks

```
template<typename Real>
static void BM_PowTemplate(benchmark::State& state) {
    std::random_device rd;
    std::mt19937 gen(rd());
    std::uniform_real_distribution<Real> dis(1, 10);
    auto s = dis(gen);
    auto t = dis(gen);
    Real y;
    while (state.KeepRunning()) {
        benchmark::DoNotOptimize(y = std::pow(s, t));
    }
    std::ostream cnull(nullptr);
    cnull << y;
}

BENCHMARK_TEMPLATE(BM_PowTemplate, float);
BENCHMARK_TEMPLATE(BM_PowTemplate, double);
BENCHMARK_TEMPLATE(BM_PowTemplate, long double);
```

Templated Benchmarks

The results are sometimes surprising; for instance double is found to be faster than float:

Run on (1 X 2300 MHz CPU)

2016-06-25 00:07:26

| Benchmark | Time | | CPU Iterations |
|-----------------------------|--------|--------|----------------|
| ----- | | | |
| BM_PowTemplate<float> | 136 ns | 127 ns | 5468750 |
| BM_PowTemplate<double> | 95 ns | 94 ns | 7000000 |
| BM_PowTemplate<long double> | 404 ns | 403 ns | 1699029 |

View algorithm scaling

Sometimes you need to analysis the scaling properties of your algorithm. Let's try an example with an algorithm with terrible scaling: Recursive Fibonnaci numbers:

```
uint64_t fibr(uint64_t n)
{
    if (n == 0)
        return 0;

    if (n == 1)
        return 1;

    return fibr(n-1)+fibr(n-2);
}
```

View algorithm scaling

Our benchmark code is

```
static void BM_FibRecursive(benchmark::State& state)
{
    uint64_t y;
    while (state.KeepRunning())
    {
        benchmark::DoNotOptimize(y = fib1(state.range_x()));
    }
    std::ostream cnull(nullptr);
    cnull << y;
}
BENCHMARK(BM_FibRecursive)->RangeMultiplier(2)->Range(1, 1<<5);
```

View algorithm scaling

BENCHMARK(BM_FibRecursive)->RangeMultiplier(2)->Range(1, 1<<5)

will request a benchmark with `state.range_x()` taking values of [1, 2, 4, 8, 16, 32].

View algorithm scaling

Run on (1 X 2300 MHz CPU)

2016-06-25 00:38:14

| Benchmark | Time | | CPU Iterations | | |
|--------------------|----------|----|----------------|----|----------|
| ----- | | | | | |
| BM_FibRecursive/1 | 7 | ns | 7 | ns | 83333333 |
| BM_FibRecursive/2 | 15 | ns | 15 | ns | 72916667 |
| BM_FibRecursive/4 | 37 | ns | 37 | ns | 17156863 |
| BM_FibRecursive/8 | 268 | ns | 268 | ns | 2868852 |
| BM_FibRecursive/16 | 13420 | ns | 13392 | ns | 64815 |
| BM_FibRecursive/32 | 24372253 | ns | 24320000 | ns | 25 |

View algorithm scaling

Can we empirically determine the asymptotic complexity of the recursive Fibonacci number calculation?

Pretty much . . . !

View algorithm scaling

google/benchmark will try to figure out the algorithmic scaling, if you ask it to. Example:

```
static void BM_FibRecursive(benchmark::State& state)
{
    uint64_t y;
    while (state.KeepRunning())
    {
        benchmark::DoNotOptimize(y = fib1(state.range_x()));
    }
    std::ostream cnull(nullptr);
    cnull << y;
    state.SetComplexityN(state.range_x());
}

BENCHMARK(BM_FibRecursive)->RangeMultiplier(2)->Range(1, 1<<5)->Complexity();
```

The result?

| | | | |
|----------------------|-------------|-------------|----------|
| BM_FibRecursive/1 | 9 ns | 9 ns | 72916667 |
| BM_FibRecursive/2 | 19 ns | 19 ns | 44871795 |
| BM_FibRecursive/4 | 42 ns | 43 ns | 14112903 |
| BM_FibRecursive/8 | 270 ns | 268 ns | 2611940 |
| BM_FibRecursive/16 | 11305 ns | 11264 ns | 54687 |
| BM_FibRecursive/32 | 27569038 ns | 27555556 ns | 27 |
| BM_FibRecursive_BigO | 828.24 N^3 | 827.83 N^3 | |
| BM_FibRecursive_RMS | 31 % | 31 % | |

It erroneously labels the algorithm as cubic; though we can see the fit is not tight.

Passing a lambda to Complexity()

If google/benchmark only tries to fit to the most common complexity classes. You are free to specify the asymptotic complexity yourself, and have google/benchmark determine goodness of fit.

The complexity of the recursive Fibonacci algorithm is ϕ^n , where $\phi := (1 + \sqrt{5})/2$ is the golden ratio, so let's use ϕ^n as a lambda . . .

Passing a lambda to Complexity()

Syntax:

```
BENCHMARK(BM_FibRecursive)
```

```
    ->RangeMultiplier(2)
```

```
        ->Range(1, 1<<5)
```

```
            ->Complexity([](int n) {return std::pow((1+std::sqrt(5))/2, n);});
```

Result:

| | | | |
|----------------------|-------------|-------------|----------|
| BM_FibRecursive/1 | 9 ns | 9 ns | 79545455 |
| BM_FibRecursive/2 | 19 ns | 19 ns | 44871795 |
| BM_FibRecursive/4 | 37 ns | 37 ns | 20588235 |
| BM_FibRecursive/8 | 228 ns | 228 ns | 3125000 |
| BM_FibRecursive/16 | 9944 ns | 9943 ns | 60345 |
| BM_FibRecursive/32 | 23910141 ns | 23866667 ns | 30 |
| BM_FibRecursive_BigO | 4.91 f(N) | 4.90 f(N) | |
| BM_FibRecursive_RMS | 0 % | 0 % | |

(Note: All lambdas are denoted as $f(N)$, so you have to know what you wrote in your source code to understand the scaling.)

Other tricks:

Standard complexity classes don't need to be passed as lambdas:

```
Complexity(benchmark::o1);
```

```
Complexity(benchmark::oLogN);
```

```
Complexity(benchmark::oN);
```

```
Complexity(benchmark::oNLogN);
```

```
Complexity(benchmark::oNSquared);
```

```
Complexity(benchmark::oNCubed);
```

Benchmarking data transfer

To get data about data transfer rates, use `state.SetBytesProcessed`¹:

```
static void BM_memcpy(benchmark::State& state) {  
    char* src = new char[state.range_x()]; char* dst = new char[state.range_x()];  
    memset(src, 'x', state.range_x());  
    while (state.KeepRunning()) {  
        memcpy(dst, src, state.range_x());  
    }  
    state.SetBytesProcessed(int64_t(state.iterations()) *  
                           int64_t(state.range_x()));  
    delete[] src; delete[] dst;  
}  
BENCHMARK(BM_memcpy)->Range(8, 8<<10);
```

¹ Taken directly from the [google/benchmark](https://github.com/google/benchmark) docs.

Benchmarking data transfer

state.SetBytesProcessed adds another column to the output:

Run on (1 X 2300 MHz CPU)

2016-06-25 19:25:23

| Benchmark | Time | CPU Iterations | | |
|----------------|--------|----------------|----------|-------------|
| ----- | | | | |
| BM_memcpy/8 | 9 ns | 9 ns | 87500000 | 883.032MB/s |
| BM_memcpy/16 | 9 ns | 9 ns | 79545455 | 1.70305GB/s |
| BM_memcpy/32 | 9 ns | 8 ns | 79545455 | 3.50686GB/s |
| BM_memcpy/64 | 11 ns | 11 ns | 62500000 | 5.35243GB/s |
| BM_memcpy/128 | 13 ns | 13 ns | 53030303 | 8.97969GB/s |
| BM_memcpy/256 | 16 ns | 16 ns | 44871795 | 15.1964GB/s |
| BM_memcpy/512 | 19 ns | 20 ns | 36458333 | 24.4167GB/s |
| BM_memcpy/1024 | 25 ns | 25 ns | 28225806 | 38.6756GB/s |
| BM_memcpy/2k | 50 ns | 50 ns | 10000000 | 38.147GB/s |
| BM_memcpy/4k | 122 ns | 122 ns | 5833333 | 31.2534GB/s |
| BM_memcpy/8k | 220 ns | 220 ns | 3240741 | 34.726GB/s |

Running a subset of your benchmarks:

Each benchmark takes ~1 second to run. If you need to analyze only one of the benchmarks, use the `--benchmark_filter` option:

```
./run_benchmarks.x --benchmark_filter=BM_memcpy/32
```

Run on (1 X 2300 MHz CPU)

2016-06-25 19:34:24

| Benchmark | Time | CPU | Iterations | |
|---------------|---------|---------|------------|-------------|
| ----- | | | | |
| BM_memcpy/32 | 11 ns | 11 ns | 79545455 | 2.76944GB/s |
| BM_memcpy/32k | 2181 ns | 2185 ns | 324074 | 13.9689GB/s |
| BM_memcpy/32 | 12 ns | 12 ns | 54687500 | 2.46942GB/s |
| BM_memcpy/32k | 1834 ns | 1837 ns | 357143 | 16.6145GB/s |

Error bars

In general you will get a pretty good idea about the standard deviation of the measurements just by running it a few times. However, if you want error bars, just specify the number of times you want your benchmark repeated:

```
BENCHMARK(BM_Pow)->Repetitions(12);
```

Error bars

Run on (1 X 2300 MHz CPU)

2016-06-25 19:57:40

| Benchmark | Time | CPU | Iterations |
|--------------------------|-------|-------|------------|
| ----- | | | |
| BM_Pow/repeats:12 | 70 ns | 70 ns | 10294118 |
| BM_Pow/repeats:12 | 73 ns | 73 ns | 10294118 |
| BM_Pow/repeats:12 | 71 ns | 71 ns | 10294118 |
| BM_Pow/repeats:12 | 70 ns | 70 ns | 10294118 |
| BM_Pow/repeats:12 | 71 ns | 71 ns | 10294118 |
| BM_Pow/repeats:12 | 72 ns | 72 ns | 10294118 |
| BM_Pow/repeats:12 | 73 ns | 73 ns | 10294118 |
| BM_Pow/repeats:12 | 71 ns | 71 ns | 10294118 |
| BM_Pow/repeats:12 | 70 ns | 70 ns | 10294118 |
| BM_Pow/repeats:12 | 73 ns | 73 ns | 10294118 |
| BM_Pow/repeats:12 | 71 ns | 71 ns | 10294118 |
| BM_Pow/repeats:12 | 76 ns | 75 ns | 10294118 |
| BM_Pow/repeats:12_mean | 72 ns | 72 ns | 10294118 |
| BM_Pow/repeats:12_stddev | 2 ns | 2 ns | 0 |